

T71 - Evaluation Matrices

March 6, 2021

1 Evaluation matrices (binary classification)

1.1 Setting up

- Breast cancer data
- 2 classes
- 30 features
- SVC

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline

# Breast cancer data
from sklearn.datasets import load_breast_cancer

# Load data
dataObj = load_breast_cancer()
X = dataObj.data
y = dataObj.target
print(np.unique(y))
print(X.shape)

# Splitting data
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.20,
                                                    stratify=y,
                                                    random_state=1)

# Constructing a pipeline object
pipe_svc = Pipeline([('scl', StandardScaler()),
                      ('clf', SVC(random_state=1))])
```

1.2 Confusion matrix

```
[ ]: from sklearn.metrics import confusion_matrix

# Training
pipe_svc.fit(X_train, y_train)

# Prediction from test data
y_pred = pipe_svc.predict(X_test)

# Confusion matrix
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
print(confmat)
```

```
[ ]: from sklearn.metrics import plot_confusion_matrix

plot_confusion_matrix(estimator=pipe_svc, X=X_test, y_true=y_test)
plt.show()
```

```
[ ]: # Note that the class 0 is "malignant" and class 1 is "benign".
print(dataObj.target_names)
```

1.2.1 Note

- The class 0 samples that are correctly predicted as class 0 are now in the upper left corner of the matrix.
- In order to change the ordering, we can use the “labels” argument.

```
[ ]: confmat = confusion_matrix(y_true=y_test, y_pred=y_pred, labels=[1, 0])
print(confmat)
```

```
[ ]: plot_confusion_matrix(estimator=pipe_svc, X=X_test, y_true=y_test, labels=[1,0])
plt.show()
```

1.3 Accuracy, Precision, Recall, and F1

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Accuracy
ACC = accuracy_score(y_true=y_test, y_pred=y_pred)
print(f"Accuracy:{ACC:6.3f}")

# Precision
PRE = precision_score(y_true=y_test, y_pred=y_pred)
print(f"Precision:{PRE:6.3f}")

# Recall
```

```

REC = recall_score(y_true=y_test, y_pred=y_pred)
print(f"Recall:{REC:6.3f}")

# F1
F1 = f1_score(y_true=y_test, y_pred=y_pred)
print(f"F1:{REC:6.3f}")

```

1.4 Use F1 score in grid search

```

[ ]: from sklearn.metrics import make_scorer

# Making score.
# Note that by default the positive class is 1. We can set the positive class_
    ↳ to 0.
scorer = make_scorer(f1_score, pos_label=0)

```

```

[ ]: from sklearn.model_selection import GridSearchCV

c_gamma_range = [0.01, 0.1, 1.0, 10.0]
param_range = [0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0]
set1 = {'clf__C': param_range, 'clf__kernel': ['linear']}
set2 = {'clf__C': param_range, 'clf__gamma': param_range, 'clf__kernel':_
    ↳ ['rbf']}
param_grid = [set1, set2]

# Grid search. Note the "scoring" argument
gs = GridSearchCV(estimator=pipe_svc,
                  param_grid=param_grid,
                  scoring=scorer,
                  cv=10,
                  n_jobs=-1)

gs = gs.fit(X_train, y_train)
print(gs.best_score_)
print(gs.best_params_)

```